



RAM SHANKAR

Senior Cloud Data & AI Architect

Phone: +91 9080968625 | Email: ram.shankar.baskaran@gmail.com | LinkedIn:

<https://www.linkedin.com/in/ram-shankar-5b9519200/> | Portfolio:

<https://ramshankar11.github.io/my-portfolio/> | Location: India (Open to international relocation)

PROFESSIONAL SUMMARY

Highly accomplished Senior Cloud Data & AI Architect with 5+ years of experience, specializing in architecting petabyte-scale data platforms and advanced Generative AI solutions for global enterprises. Proven expertise in cloud modernization across GCP, AWS, and Azure, driving digital transformation and cost optimization through FinOps and robust CI/CD pipelines. A strategic leader with a strong commitment to performance engineering, dedicated to delivering data-driven innovations and fostering team collaboration.

WORK EXPERIENCE

LatentView Analytics | Senior Analyst - Data Engineering

Jan 2026 - Present

- Architected and implemented a multi-layered Data Lakehouse on GCP BigQuery, engineering robust ETL pipelines via Cloud Composer to process 5TB+ of daily customer lead data, significantly enhancing targeting capabilities and conversion metrics for global sales teams.
- Spearheaded the design and deployment of an enterprise-grade Retrieval-Augmented Generation (RAG) system using Vertex AI and LangChain, empowering stakeholders with instant access to internal knowledge bases and reducing data retrieval time by 40%.
- Optimized DevOps workflows by standardizing infrastructure provisioning with modular Terraform scripts, cutting environment setup time from 3 days to 2 hours across all environments.
- Provided technical leadership and mentorship to a team of 4 engineers on Kubernetes (GKE) deployment strategies and Python best practices, elevating team code quality scores by 30%.

STG Infotech | Software Engineer

Apr 2024 - Jan 2026

- Led the end-to-end migration of 50+ complex legacy Alteryx workflows to a highly scalable, cloud-native architecture utilizing Dataflow (Apache Beam) and BigQuery, achieving annual licensing cost savings of \$45,000.
- Developed innovative Agentic AI workflows leveraging Large Language Models (LLMs) to automate data quality incident reporting and Jira ticket assignment, reducing Mean Time to Resolution (MTTR) by 25%.
- Engineered a high-throughput real-time event processing pipeline with GCP Pub/Sub and Dataflow for IoT sensor data ingestion, enabling sub-second latency for critical operational intelligence dashboards.
- Enhanced data governance and integrity by implementing Dataform for SQL transformation management, integrating CI/CD version control and automated testing.

Accenture | Data Engineering Analyst

Jun 2021 - Apr 2024

- Optimized and maintained robust ETL pipelines using Python and SQL, applying FinOps principles to reduce BigQuery slot usage and achieve a 20% reduction in overall query costs.
- Ensured high platform reliability and 99.9% uptime for mission-critical enterprise data workflows by expertly managing Google Cloud Composer (Airflow) environment upgrades and complex dependency management.
- Developed and maintained robust CI/CD pipelines using GitHub Actions, automating the deployment of

Dockerized applications to Google Kubernetes Engine (GKE) and eliminating manual deployment errors.

TECHNICAL SKILLS

Languages & Core: Python (Advanced), SQL (Advanced), Bash, Go, PostgreSQL

Cloud & Infrastructure: Google Cloud Platform (BigQuery, Dataflow, Pub/Sub, Vertex AI, Cloud Run), AWS (Lambda, S3, DynamoDB), Azure, Terraform, Docker, Kubernetes (GKE, EKS), Helm, GitHub Actions, ArgoCD, Prometheus

Data Technologies & AI: Apache Airflow (Cloud Composer), dbt (Data Build Tool), Apache Kafka, PySpark, Alteryx, Databricks, RAG Pipelines, LangChain, Vector Databases (Pinecone, Weaviate), Agentic AI, OpenAI API, Gemini Pro, Great Expectations, Dataform

CORE COMPETENCIES & DOMAIN

Cloud Architecture & Strategy | Data Engineering & Analytics | Generative AI & Machine Learning | DevOps & SRE | FinOps & Cost Optimization | Data Governance & Quality | Digital Transformation | Technical Leadership | Agile Methodologies | Performance Engineering

EDUCATION

Bachelor of Engineering (B.E.), Computer Science | Tamil Nadu, India

P.S.R. Engineering College | 2021 | 7.9 GPA

KEY PROJECTS

Automotive Supply Chain Migration (Legacy Modernization) | Work

Tech Stack: GCP (Dataflow, BigQuery, Cloud Composer), Alteryx

- Led the migration of mission-critical supply chain logic from on-premise Alteryx to a serverless GCP architecture, reducing operational costs by 40% and data latency from 24 hours to 1 hour.

Insurance Policy RAG Agent (Generative AI) | Work

Tech Stack: Vertex AI, LangChain, Vector Search

- Designed an internal Agentic AI tool for insurance underwriters using Vertex AI, LangChain, and Vector Search, streamlining underwriting and reducing manual document search time by 60%.

Enterprise GenAI Assistant & Agentic Process Automation | Work

Tech Stack: GenAI Platform, RAG, BigQuery SQL, Playwright

- Architected a unified GenAI platform to automate supply chain operations, featuring a RAG Knowledge Engine, Transactional Agents (Tool Calling to BigQuery SQL), and Autonomous Browser Agents (Playwright) for vendor invoice processing.

Real-Time Streaming & Data Quality Platform | Personal

Tech Stack: Kafka, Python, Docker, Great Expectations

- Engineered an event-driven log processing and anomaly detection architecture with Kafka, Python, and Docker, integrating Great Expectations for automated schema validation and high data fidelity.

Modern Analytics Engineering Stack | Personal

Tech Stack: dbt, PostgreSQL, SQL

- Built an end-to-end batch ETL pipeline for Crypto market data into a PostgreSQL warehouse, leveraging dbt and advanced SQL for layered dimensional modeling and analytics engineering.

Enterprise GCP Infrastructure Automation | Personal

Tech Stack: Terraform, GCP Compute, Kubernetes (GKE), Helm, Cloud Run, Pub/Sub

- Provisioned a secure 3-Tier cloud architecture using Terraform and GCP Compute, deploying scalable microservices with Kubernetes (GKE), Helm, Cloud Run, and Pub/Sub.

Multi-Cloud Serverless Data Lake & API | Personal

Tech Stack: AWS Lambda, S3, LocalStack, AWS SAM, DynamoDB, Azure Bicep, Ansible

- Architected an event-triggered file processing pipeline using AWS Lambda, S3, and LocalStack, developed highly available REST APIs with AWS SAM and DynamoDB, and implemented hybrid infrastructure using Azure Bicep and Ansible.

CERTIFICATIONS

- Google Cloud: Professional Cloud Architect
- Google Cloud: Professional Data Engineer
- Google Cloud: Professional Machine Learning Engineer
- AWS: Certified Solutions Architect - Associate
- Microsoft Azure: Data Scientist Associate (DP-100)
- Red Hat: Certified System Administrator (RHCSA)
- GitHub: GitHub Actions
- GitHub: GitHub Foundations